



Checked By	Dr. Avinash Golande
Sign	
Date	
Grade	

EXPERIMENT NO . 03

- **Title:-** Study and Implementation of *libvirt* for Virtual Machine (VM) Management
- **Objective :-**
 - To understand libvirt as a virtualization management tool.
 - To learn how to create, configure, and manage VMs using libvirt.
 - To explore Python bindings of libvirt for programmatically controlling hypervisors.
- **Resources used :-** PC / Laptop with Linux OS (Ubuntu/Debian recommended) ,libvirt-daemon, qemu-kvm, virt-manager , Python 3 with libvirt-python module installed
- **Theory :-**

What is libvirt?

libvirt is basically the “**control tower**” for managing virtual machines, containers, and virtualization resources on a host system. Instead of directly talking to a specific hypervisor like QEMU/KVM, Xen, VirtualBox, or VMware ESXi, we interact with **libvirt**, which then translates our requests into the correct commands.

It includes:

- **Library:** libvirt.so
- **Daemon/Service:** libvirtd (or modular daemons in newer versions)
- **Tools:** virsh, virt-install, virt-manager

Main Components:

- I. **libvirt API / library** – Interfaces for programming in C, Python, Go, Java.
- II. **libvirtd service** – Handles communication with hypervisors (locally or remotely).
- III. **CLI tools:**
 - virsh – Command line VM control
 - virt-install – For creating VMs

- virt-manager – GUI frontend

Capabilities of libvirt:

- Create and delete VMs
- Configure CPU, RAM, disks, and networks
- Attach/detach resources dynamically
- Live migrate VMs
- Take snapshots (disk + memory)
- Manage storage pools and virtual networks

- **Option A — Real libvirt + Python bindings (create / start / stop VMs)**

Prerequisites

- I. Linux host with libvирtd running and qemu/kvm installed.
- II. Python 3 and libvirt Python package:

bash

CopyEdit

```
sudo apt-get install -y libvirt-daemon-system qemu-kvm virtinst # Debian/Ubuntu example
pip install libvirt-python
```

- I. qemu-img available to create QCOW2 images.
 1. You run the script as a user with permission to talk to libvirt (often your user needs to be in libvirt/libvirt-qemu group or run via sudo).

Note: this script issues shell commands like qemu-img and manipulates domain XML — be careful on production systems.

- I. What the script does
- II. Create a backing QCOW2 image.
- III. Create a minimal cloud-init ISO (optional) or use injected cloud-init user-data.
- IV. Define a domain from an XML template and start the VM.
- V. Provide functions to list, shutdown, destroy, undefine, snapshot.

```
#!/usr/bin/env python3
```

```
"""
```

Simple libvirt VM manager using libvirt Python bindings.

Provides: create_disk, define_domain, start, shutdown, destroy, undefine, list_domains

"""

```
import libvirt
import subprocess
import os
import tempfile
import time
from xml.etree import ElementTree as ET
```

```
LIBVIRT_URI = "qemu:///system" # local system
```

```
def qemu_img_create(path, size_gb=10, backing=None):
    """Create a qcow2 image. If backing is provided, create a copy-on-write image."""
    cmd = ["qemu-img", "create", "-f", "qcow2"]
    if backing:
        cmd += ["-b", backing]
    cmd += [path, f"{size_gb}G"]
    subprocess.check_call(cmd)
    return path
```

```
def create_cloud_init_iso(hostname, ssh_authorized_key, iso_path):
    """Create a minimal cloud-init ISO with user-data/meta-data (for cloud images)."""
    tmpdir = tempfile.mkdtemp()
    user_data = f"""#cloud-config
hostname: {hostname}
ssh_authorized_keys:
- {ssh_authorized_key}
"""

    with open(os.path.join(tmpdir, "user-data"), "w") as f:
```

```

f.write(user_data)

with open(os.path.join(tmpdir, "meta-data"), "w") as f:
    f.write("instance-id: iid-local01\nlocal-hostname: {}\n".format(hostname))

# genisoimage or mkisofs

cmd = ["genisoimage", "-output", iso_path, "-volid", "cidata", "-joliet", "-rock",
        os.path.join(tmpdir, "user-data"), os.path.join(tmpdir, "meta-data")]
subprocess.check_call(cmd)

return iso_path

def get_conn(uri=LIBVIRT_URI):
    return libvirt.open(uri)

def generate_domain_xml(name, disk_path, iso_path=None,
                       memory_mb=2048, vcpus=2, mac="52:54:00:12:34:56"):
    """
    Generate a basic domain XML for qemu+kvm.

    Adjust according to your environment (network, disk bus, etc.)
    """

    xml = f"""
<domain type='kvm'>
    <name>{name}</name>
    <memory unit='MiB'>{memory_mb}</memory>
    <vcpu placement='static'>{vcpus}</vcpu>
    <os>
        <type arch='x86_64' machine='pc'>hvm</type>
        <boot dev='hd'/>
    </os>
    <features>
        <acpi/>
        <apic/>
    </features>
</domain>
"""

    if iso_path:
        xml += f'    <disk type="file" device="cdrom">
        <source file="{iso_path}" />
        <target dev="cdrom" />
    </disk>'

    if disk_path:
        xml += f'    <disk type="file" device="disk">
        <source file="{disk_path}" />
        <target dev="hda" />
        <backing-store type="file" />
    </disk>'

    return xml

```

```

<pae/>
</features>
<clock offset='utc' />
<on_poweroff>destroy</on_poweroff>
<on_reboot>restart</on_reboot>
<on_crash>destroy</on_crash>
<devices>
    <emulator>/usr/bin/qemu-system-x86_64</emulator>
    <disk type='file' device='disk' />
        <driver name='qemu' type='qcow2' />
        <source file='{disk_path}' />
        <target dev='vda' bus='virtio' />
    </disk>
    """
    if iso_path:
        xml += f"""
        <disk type='file' device='cdrom' />
            <driver name='qemu' type='raw' />
            <source file='{iso_path}' />
            <target dev='hdc' bus='ide' />
            <readonly />
        </disk>
        """
        xml += f"""
        <interface type='network' />
            <mac address='{mac}' />
            <source network='default' />
            <model type='virtio' />
        </interface>
        <serial type='pty' ><target port='0' /></serial>

```

```

<console type='pty'><target type='serial' port='0'></console>
<input type='tablet' bus='usb'>
<graphics type='vnc' port='-1' listen='0.0.0.0' />
<memballoon model='virtio' />
</devices>
</domain>
"""

return xml

def define_and_start(vm_name, disk_path, iso_path=None, memory_mb=2048, vcpus=2):
    conn = get_conn()
    xml = generate_domain_xml(vm_name, disk_path, iso_path, memory_mb, vcpus)
    try:
        dom = conn.defineXML(xml)
    except libvirt.libvirtError as e:
        print("Failed to define domain:", e)
        conn.close()
        raise
    print(f"Defined domain {vm_name}")
    dom.create() # boot it
    print(f"Started domain {vm_name}")
    conn.close()
    return dom

def list_domains():
    conn = get_conn()
    ids = conn.listAllDomains()
    out = []
    for dom in ids:
        state, reason = dom.state()

```

```

        out.append((dom.name(), dom.ID(), state))

    conn.close()

    return out


def shutdown_domain(name):
    conn = get_conn()
    dom = conn.lookupByName(name)
    dom.shutdown()
    conn.close()


def destroy_and_undefine(name):
    conn = get_conn()
    dom = conn.lookupByName(name)
    if dom.isActive():
        dom.destroy()
    dom.undefine()
    conn.close()


# Example usage if run as script
if __name__ == "__main__":
    # PARAMETERS - change as needed
    vm_name = "demo-vm-1"
    disk_file = f"/var/lib/libvirt/images/{vm_name}.qcow2"
    base_image = "/var/lib/libvirt/images/ubuntu-22.04-server-cloudimg-amd64.img" # must exist
    iso_path = f"/var/lib/libvirt/images/{vm_name}-cidata.iso"
    ssh_pubkey = "ssh-rsa AAAA...yourkey... user@example"
    # create disk (COW) backed by base image
    if not os.path.exists(disk_file):
        print("Creating qcow2 disk...")
        qemu_img_create(disk_file, size_gb=10, backing=base_image)

```

```

# create cloud-init ISO

if not os.path.exists(iso_path):
    create_cloud_init_iso(vm_name, ssh_pubkey, iso_path)

# define and start

define_and_start(vm_name, disk_file, iso_path, memory_mb=2048, vcpus=2)
print("VMs:", list_domains())

```

- Option B — VM Scheduling Simulator (pure Python)

If you want to *simulate scheduling* (pack VMs onto hosts using heuristics), here's a portable simulator you can run locally. It models hosts and VMs (CPU cores and RAM) and provides two placement heuristics: First Fit and Best Fit by remaining capacity.

Save as `vm_scheduler_sim.py` and run with Python 3.

```
#!/usr/bin/env python3
```

```
"""
```

VM scheduling simulator.

Hosts have capacities: `cpu_cores`, `ram_gb`

VMs request: `cpu`, `ram`

Heuristics: first-fit, best-fit (by residual RAM)

```
"""
```

```
from dataclasses import dataclass, field
```

```
from typing import List, Tuple
```

```
import random
```

```
import pprint
```

```
@dataclass
```

```
class VM:
```

```
    id: str
```

```
    cpu: int
```

```
    ram: int # GB
```

```

@dataclass
class Host:
    id: str
    total_cpu: int
    total_ram: int
    used_cpu: int = 0
    used_ram: int = 0
    vms: List[VM] = field(default_factory=list)

    @property
    def free_cpu(self):
        return self.total_cpu - self.used_cpu

    @property
    def free_ram(self):
        return self.total_ram - self.used_ram

    def can_host(self, vm: VM) -> bool:
        return (self.free_cpu >= vm.cpu) and (self.free_ram >= vm.ram)

    def place_vm(self, vm: VM):
        if not self.can_host(vm):
            raise ValueError(f'Host {self.id} cannot host VM {vm.id}')
        self.used_cpu += vm.cpu
        self.used_ram += vm.ram
        self.vms.append(vm)

    def schedule_first_fit(hosts: List[Host], vms: List[VM]) -> Tuple[List[Host], List[VM]]:
        unplaced = []
        for vm in vms:

```

```
placed = False
```

```
for host in hosts:
```

```
    if host.can_host(vm):
```

```
        host.place_vm(vm)
```

```
        placed = True
```

```
        break
```

```
if not placed:
```

```
    unplaced.append(vm)
```

```
return hosts, unplaced
```

```
def schedule_best_fit_ram(hosts: List[Host], vms: List[VM]) -> Tuple[List[Host], List[VM]]:
```

```
    unplaced = []
```

```
    for vm in vms:
```

```
        # find host with smallest free_ram that can still fit vm (best-fit)
```

```
        candidates = [h for h in hosts if h.can_host(vm)]
```

```
        if not candidates:
```

```
            unplaced.append(vm)
```

```
            continue
```

```
        best = min(candidates, key=lambda h: h.free_ram - vm.ram)
```

```
        best.place_vm(vm)
```

```
    return hosts, unplaced
```

```
def random_test():
```

```
    hosts = [Host(f'host{i}', total_cpu=32, total_ram=128) for i in range(4)]
```

```
    # create a set of VMs with random sizes
```

```
    vms = []
```

```
    for i in range(20):
```

```
        cpu = random.choice([1,2,4,8])
```

```
        ram = random.choice([1,2,4,8,16])
```

```
        vms.append(VM(f'vm{i}', cpu, ram))
```

```

print("Scheduling with First-Fit")

import copy

h1 = [copy.deepcopy(h) for h in hosts]

schedule_first_fit(h1, vms)

for h in h1:

    print(h.id, "used_cpu", h.used_cpu, "used_ram", h.used_ram, "vms", len(h.vms))

print()

print("Scheduling with Best-Fit (RAM)")

h2 = [copy.deepcopy(h) for h in hosts]

schedule_best_fit_ram(h2, vms)

for h in h2:

    print(h.id, "used_cpu", h.used_cpu, "used_ram", h.used_ram, "vms", len(h.vms))

if __name__ == "__main__":
    random_test()

```

- Example workflow: use both together
 - I. Design VM descriptors (CPU, RAM, image, cloud-init key, name).
 - II. Run simulator to decide which host each VM should land on (use the simulator heuristics).
 - III. Write a small orchestration script that takes the simulator's placements and uses the libvirt script (Option A) to connect to each host's libvirt URI and define/start the VM there:
 - For remote hosts: qemu+ssh://user@host/system can be used as LIBVIRT_URI in Option A.
 - a. Use SSH keys for passwordless remote libvirt access.

Conclusion:

Successfully studied and implemented **libvirt** as a virtualization management framework.

- Understood its architecture (library, daemon, CLI tools).
- Learned real VM management with libvirt Python bindings.
- Explored VM scheduling simulator to test placement strategies.